

Uncertainty in Recurrent Decision Tree Classifiers

Stefan Wezel

Explainable Machine Learning

October 19, 2020

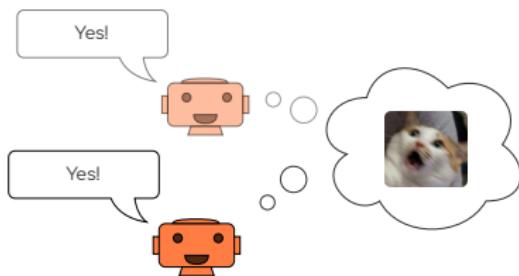
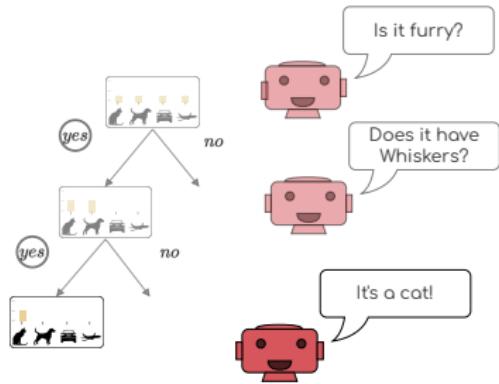
What?

Setting

- There are a lot of architectures that perform great on image classification tasks
- Maybe, most prominently: ResNet
- However, they only yield a classification
- In many settings a classification is not worth much without the reasoning behind it

What?

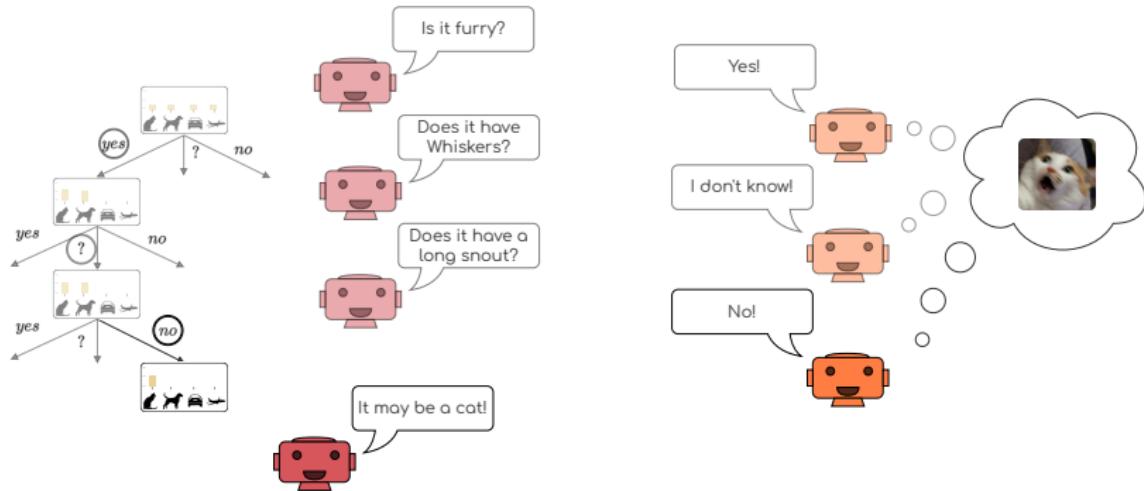
Recap



- Two agents
- One is asking questions and one is answering them
- The unfolding decision process is an interpretable tree

What?

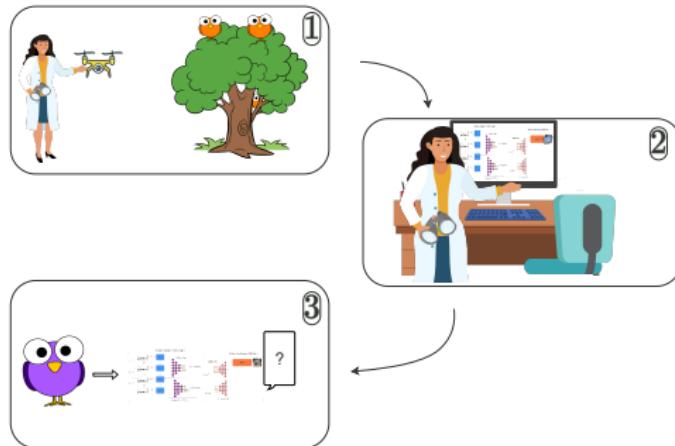
Recap



- Two agents
- One is asking questions and one is answering them
- The unfolding decision process is an interpretable tree

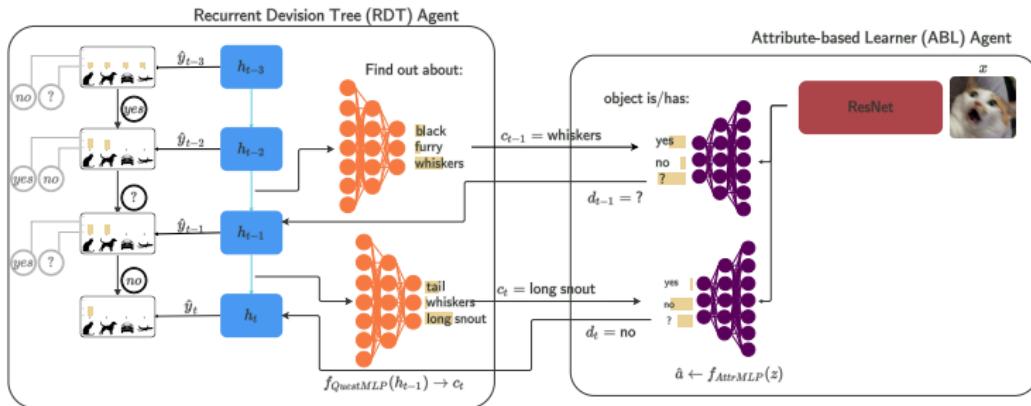
Why do we need uncertainty?

A Practical Example...



- The ornithologist is tasked to survey bird species, which she automates using a drone and computer vision software
- She uses our model to go through the vast amount of collected data
- Some bird species unknown to the model appear in the data. The model yields high uncertainty and the ornithologist can classify them manually

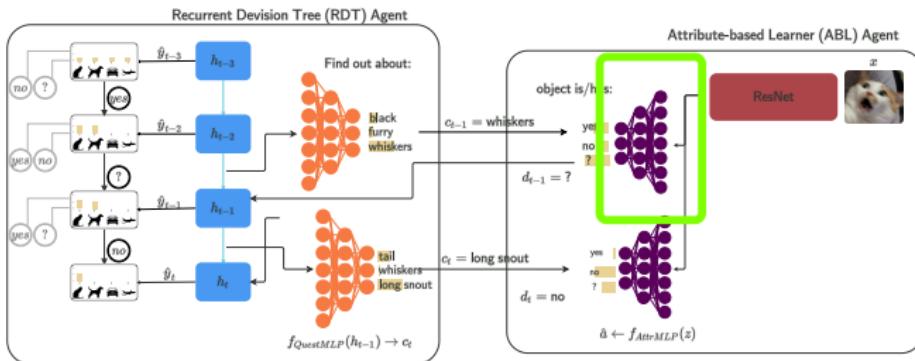
How? Architecture



- The RDT can not see the images but only ask whether an attribute is present in the image
- The AbL can see the image and based on features answer the RDT's questions

How?

Do we get uncertainty information?



- We use dropout uncertainty estimation to retrieve uncertainty (more on that later)
- We want the AbL to say 'I don't know' in the case of high uncertainty
- However, we want to keep the ResNet as feature extractor
- Thus, we use dropout in $f_{AttrMLP}$ (more on that later)
- After extracting features, we do n forward passes and compute variance

Attribute-based Learner

Answering questions

- We extract features from a given image using a ResNet
- A MLP maps extracted features to 'yes-no' answers indicating absence/presence of attributes
- It returns a tensor with the shape number of attributes × decision size
- We get a discrete answer from our AbL though applying

$$\text{TempSoftmax}(\log \pi) = \frac{\exp((\log \pi_i)/\tau)}{\sum_{j=1}^K \exp((\log \pi_j)/\tau)} = d_t \text{ on } \log \pi$$

which are the logit values for either 'Yes', or 'No' per attribute.

- The TempSoftmax serves as differentiable approximation to a one argmax returning a one-hot encoding

Recurrent Decision Tree

Building a decision tree

- LSTM
 - Hidden state based on previous hidden states and new answers
- Explicit Memory
 - Stores all questions and corresponding answers
 - This is our decision tree
- $f_{QuestMLP}$
 - Find next question to ask based on LSTM's hidden state
- $f_{ClassMLP}$
 - Make classification based on LSTM's hidden state

Training the two agents

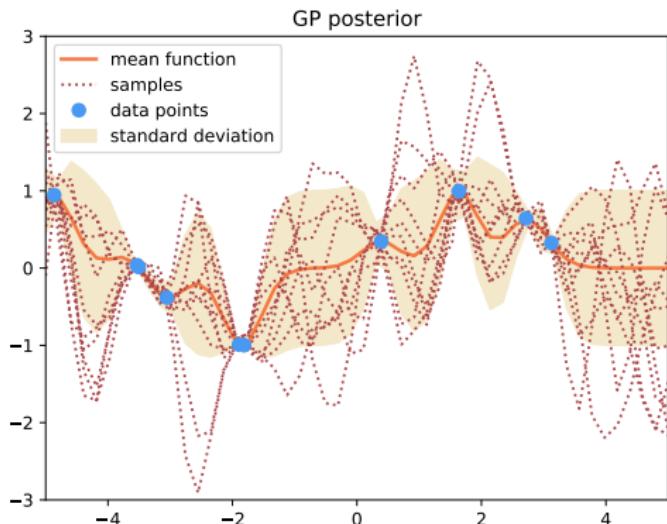
- We optimize for class and attribute accuracy
-

$$\mathcal{L} = \frac{1}{T} \sum_{t=1}^T [(1 - \lambda) \mathcal{L}_{CE}(y, \hat{y}_t) + \lambda \mathcal{L}_{CE}(\alpha_{y,c_t}, \hat{\alpha}_{c_t})]$$

- λ can be used to balance the two loss terms
- For all of our experiments, we use $\lambda = 0.2$

Background

A small excursion to Gaussian Processes (GP)



- Data points can be described by (infinitely) many functions
- A GP is a PDF over these functions
- Intuition: → a GP yields a probability for function values at any given index
- Parameterized by mean function and covariance function
- The variance resembles the model uncertainty where no data is given

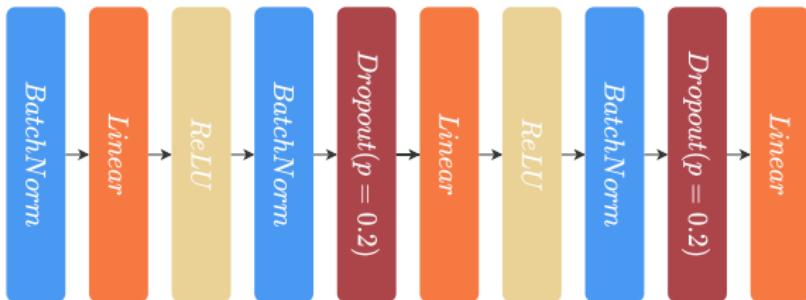
Background

Dropout Uncertainty Estimation

- A GP can be seen as an infinitely wide neural network
- In turn, we can view a neural network with a finite amount of layers as an approximation to a GP
- To compute the posterior in GP's, methods of variational inference are used
 - GP objective gets turned into a minimization objective
 - For computing covariance matrix, we use Monte-Carlo integration
- This allows us to rewrite a GP's objective as to objective of a dropout neural network

Getting Uncertainty Information

through dropout uncertainty information



- We include dropout layers in our $f_{AttrMLP}$
- We tested different configurations and a combination of batchnorm and dropout worked best

Using Uncertainty Information

as an inductive bias

- This allows us to get uncertainty information from our model
- But now, we need a way to use it
- we use two different strategies
 - We prevent the model from asking questions regarding uncertain attributes → remRDTC
 - We give the model the ability to answer with 'I don't know' as extended vocabulary → extRDTC
- For both strategies, we need to make sure the model does not use any gradients coming from uncertain attributes
- This ensures that the uncertainty information only remains an inductive bias and the model does not misuse it

Using Uncertainty Information

Removing uncertain attributes



- The output from $f_{QuestMLP}$ is an index that indicates the attribute in question
- In case, an attribute is deemed uncertain by the AbL, we replace selection logits at those indices with $-\infty$. This prevents the Gumbel softmax from picking such attributes as index

Using Uncertainty Information

Extending the vocabulary

$$\begin{array}{ccc} Y & N & ?_{initial} \\ \left[\begin{array}{cc} 1 & 0 \\ 0 & 1 \end{array} \right] & \left[\begin{array}{c} 0 \\ 0 \end{array} \right] \\ \downarrow & & \end{array}$$

$$\begin{array}{ccc} Y & N & ?_{initial} \\ \left[\begin{array}{ccc} 1 & 0 & 0 \\ 0 & 1 & 0 \end{array} \right] & \left[\begin{array}{c} 0 \\ 1 \end{array} \right] \\ \downarrow & & \end{array}$$



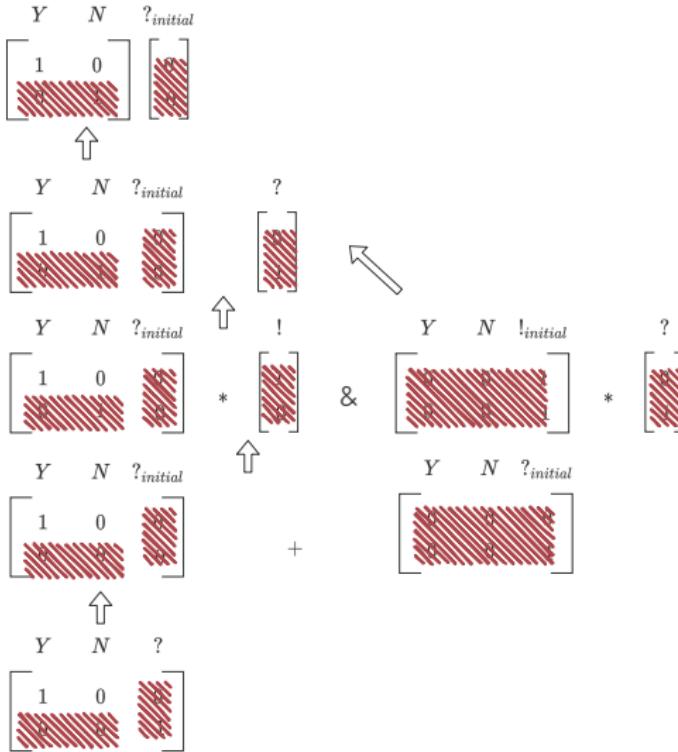
$$\begin{array}{ccc} Y & N & ?_{initial} \\ \left[\begin{array}{ccc} 1 & 0 & 0 \\ 0 & 1 & 0 \end{array} \right] & * \left[\begin{array}{c} 1 \\ 0 \end{array} \right] & \& \begin{array}{ccc} Y & N & !_{initial} \\ \left[\begin{array}{ccc} 0 & 0 & 1 \\ 0 & 0 & 1 \end{array} \right] & * \left[\begin{array}{c} 0 \\ 1 \end{array} \right] \\ \downarrow & & \end{array} \\ \& & \end{array}$$

$$\begin{array}{ccc} Y & N & ?_{initial} \\ \left[\begin{array}{ccc} 1 & 0 & 0 \\ 0 & 0 & 0 \end{array} \right] & + & \begin{array}{ccc} Y & N & ?_{initial} \\ \left[\begin{array}{ccc} 0 & 0 & 0 \\ 0 & 0 & 1 \end{array} \right] & \downarrow & \end{array} \\ \& & \end{array}$$

$$\begin{array}{ccc} Y & N & ? \\ \left[\begin{array}{ccc} 1 & 0 & 0 \\ 0 & 0 & 1 \end{array} \right] & & \end{array}$$

Using Uncertainty Information

Extending the vocabulary



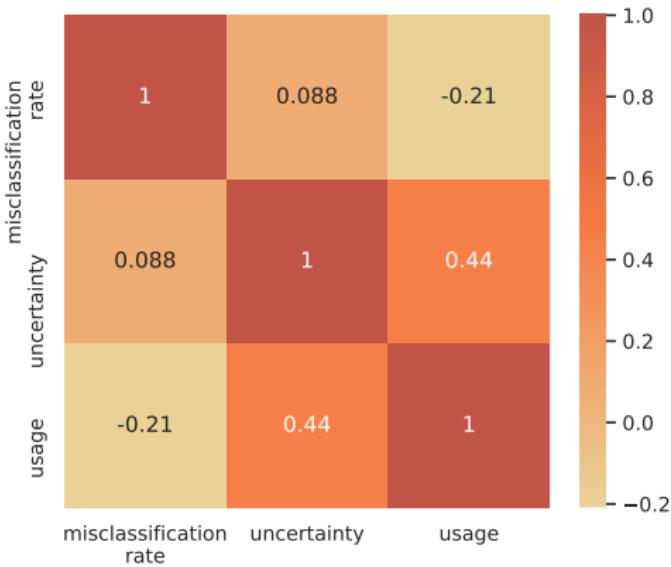
- We can interpret the variance arising from multiple forward passes in a dropout neural net as uncertainty
- We do this dropout uncertainty estimation in our $f_{AttrMLP}$
- We use the retrieved uncertainty estimate for either preventing the RDT from asking questions regarding uncertain attributes or extent the AbL's vocabulary

Experiments

- Now, we can make our model aware of, and express its uncertainties
- We use this in our experiments to:
 - Investigate uncertainty and its relationship to other variables
 - Test our model on OOD data
 - Test the model's performance on benchmark datasets

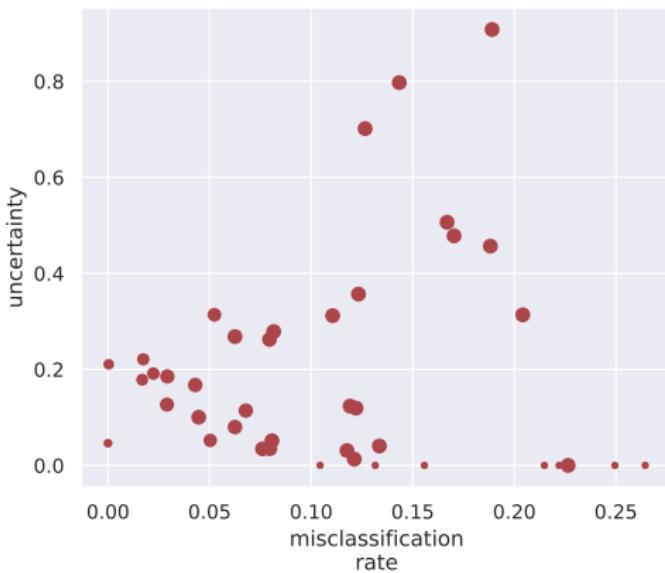
Experiments

Investigating Uncertainties



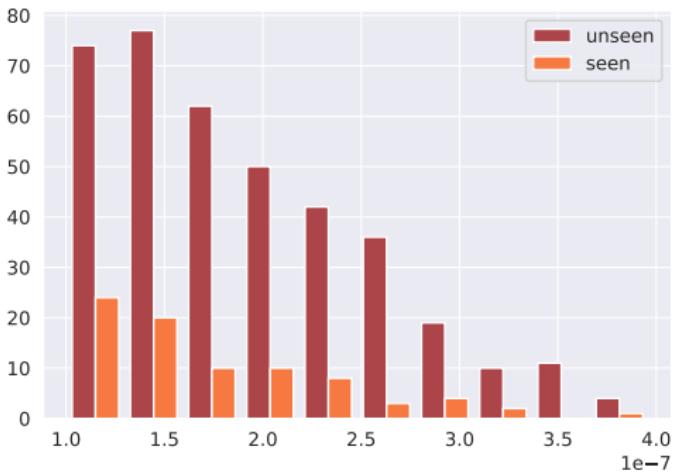
Experiments

Investigating Uncertainties



Experiments

OOD Detection



Experiments

Results on Benchmark Datasets

	AWA2	aPY	CUB
ResNet [HZRS16]	98.2 \pm 0.0	85.1 \pm 0.6	79.0 \pm 0.2
DT	78.0 \pm 0.4	64.3 \pm 0.6	19.3 \pm 0.3
dNDF[KFCRB15]	97.6 \pm 0.2	85.0 \pm 0.6	73.8 \pm 0.3
RDTC[AA19]	98.0 \pm 0.1	85.7 \pm 0.7	78.1 \pm 0.2
XDT	73.9 \pm 0.9	59.9 \pm 1.5	4.9 \pm 1.3
aRDTC[AA19]	98.6	86.1	77.9 \pm 0.6
remRDTC(ours)	98.7	86.4	77.7
extRDTC(ours)	98.7	85.4	77.8

Experiments

Results on Benchmark Datasets

	aRDTC [AA19]	Random Baseline	remRDTC	extRDTC
AWA2				
Class	98.6	98.5	98.7	98.7
Attribute	80.4	84.6	87.5	82.31
aPY				
Class	86.1	86.5	86.4	85.4
Attribute	86.4	86.2	87.6	87.12
CUB				
Class	77.9	76.8	77.7	77.8
Attribute	68.6	70.0	77.4	82.6

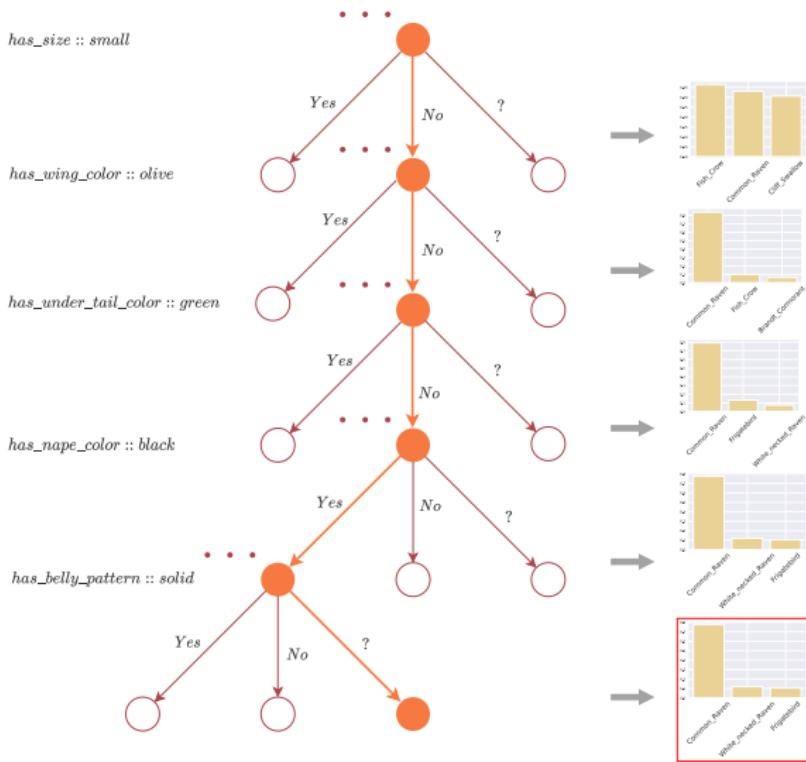
Discussion

Looking back...

- The right kind of uncertainty?
- Uncertainty versus risk
- Beyond attribute uncertainty
-

Conclusions

A qualitative Example



- [AA19] Stephan Alaniz and Zeynep Akata. Explainable observer-classifier for explainable binary decisions. arXiv preprint arXiv:1902.01780, 2019.
- [HZRS16] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In Proceedings of the IEEE conference on computer vision and pattern recognition, pages 770–778, 2016.
- [KFCRB15] Peter Kotschieder, Madalina Fiterau, Antonio Criminisi, and Samuel Rota Bulo. Deep neural decision forests. In Proceedings of the IEEE international conference on computer vision, pages 1467–1475, 2015.